

PT 1 STAGE 3 CS 411 Document

```
g563105177@cloudshell:~ (dbpjt-418422) $ gcloud sql connect dbpjt-sql --user=Qiaochu --project=dbpjt-418422
Allowlisting your IP for incoming connection for 5 minutes...done.
Connecting to database with SQL user [Qiaochu]. Enter password:
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 75305
Server version: 8.0.31-google (Google)

Copyright (c) 2000, 2024, Oracle and/or its affiliates.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> show database
-> show database;
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that corresponds to your MySQL server version for the right syntax to use near 'database
show database' at line 1
mysql> show databases;
+-----+
| Database |
+-----+
| information_schema |
| mysql |
| performance_schema |
| sys |
| weatherxwardrobe |
+-----+
5 rows in set (0.01 sec)

mysql> 
```

Table rows count:

```

mysql> SELECT COUNT(*) FROM location;
+-----+
| COUNT(*) |
+-----+
|      1009 |
+-----+
1 row in set (0.00 sec)

mysql> SELECT COUNT(*) FROM user;
+-----+
| COUNT(*) |
+-----+
|      1466 |
+-----+
1 row in set (0.00 sec)

mysql> SELECT COUNT(*) FROM weather;
+-----+
| COUNT(*) |
+-----+
|      3027 |
+-----+
1 row in set (0.00 sec)

mysql> SELECT COUNT(*) FROM apparel;
+-----+
| COUNT(*) |
+-----+
|       306 |
+-----+
1 row in set (0.00 sec)

mysql> SELECT COUNT(*) FROM usersubscribeslist;
+-----+
| COUNT(*) |
+-----+
|      3484 |
+-----+
1 row in set (0.01 sec)

```

1) Create a markdown or pdf file called "Database Design" in the doc folder of your GitHub repo.

In the Database Design markdown or pdf, provide the Data Definition Language (DDL) commands you used to create each of these tables in the database. Here's the syntax of the CREATE TABLE DDL command:

```
create table user(  
userid INT PRIMARY KEY,  
username VARCHAR(50),  
gender VARCHAR(10));
```

```
create table location (  
CityID INT PRIMARY KEY,  
CityName VARCHAR(50),  
State VARCHAR(20),  
Latitude REAL,  
Longitude REAL);
```

```
create table weather (  
date DATETIME, cityID INT PRIMARY KEY,  
Min_temp INT, Max_temp INT, Avg_temp INT,  
Wind_speed real,  
Category VARCHAR(20),  
FOREIGN KEY(cityID) REFERENCES location (cityID));
```

```
Create table apparel (  
ApparelID INT PRIMARY KEY,  
product_name VARCHAR(50),  
category VARCHAR(50),  
Temp_from INT,  
Temp_to INT  
);
```

```
create table outfit(  
outfit_id INTEGER PRIMARY KEY AUTO_INCREMENT,  
Season Varchar(10),  
Datecreated DATETIME,  
user_id INTEGER,  
apparelID INTEGER,  
FOREIGN KEY(user_id) REFERENCES user(userid),  
FOREIGN KEY(apparelID) REFERENCES apparel(ApparelId)  
);
```

```
CREATE TABLE usersubscribeslist (  
user_id INTEGER,  
date DATETIME,  
city_id INTEGER,  
FOREIGN KEY(city_id)REFERENCES weather(cityId),  
FOREIGN KEY (user_id) REFERENCES user(userid),
```

FOREIGN KEY (date) REFERENCES weather (date),
PRIMARY KEY(user_id,date, city_id))

For outfit table the records are entered using this stored procedure whenever a user requests for a new outfit for that particular weather:

Stored procedure for outfit:

```
CREATE procedure OutfitCreate(IN season VARCHAR(10), IN datecreated DATETIME, IN user_id2 INTEGER)
```

```
BEGIN
```

```
    DECLARE int outfit_id 0;
```

```
    DECLARE apparelcur CURSOR FOR select a.apparelID From apparel a  
JOIN( select w.Avg_temp FROM weather w JOIN location l ON w.cityID =l.cityID JOIN  
usersubscribeslist u ON l.cityID = u.city_ID)
```

```
AS current_weather
```

```
Where a.Temp_from <=current_weather.Avg_temp AND a.Temp_to >=
```

```
current_weather.Avg_temp where u.user_id=user_id2;
```

```
    SELECT count(*)+1 into outfit_id from outfit;
```

```
    OPEN apparelcur;
```

```
REPEAT
```

```
    FETCH apparelcur INTO apparelIDval;
```

```
    INSERT IGNORE INTO outfit VALUES(outfit_id,season,datecreated,user_id2,  
apparelIDval);
```

```
    UNTIL done
```

```
    END REPEAT;
```

```
END;
```

So it is empty at this point. We will use this procedure along with the front end to generate it whenever a user requests for it.

We had a separate table for outfit and the relationship outfit_has_apparel in stage 2.

We combined that relationship into the outfit table in the db in sql since it makes it easier to do queries on that table. Otherwise we would always do a join with outfit and outfit_has, only increasing complexity unnecessarily.

2) and 3) 4 complex queries and screenshots of top 15 results, and Indexing on the query:

a) Printing all apparel appropriate for a user's list's location's weather

```
SELECT a.category, a.product_name
FROM apparel a
JOIN (
    SELECT w.Avg_temp
    FROM weather w
    JOIN location l ON w.cityID = l.cityID
    JOIN usersubscribeslist u ON l.cityID = u.city_ID
)
AS current_weather
WHERE a.Temp_from <= current_weather.Avg_temp AND a.Temp_to >= current_weather.Avg_temp
```

category	product_name
Jacket	HOODED TECHNICAL JACKET
Jacket	HOODED TECHNICAL JACKET
Jacket	HOODED TECHNICAL JACKET
Jacket	LINEN - COTTON OVERSHIRT
Jacket	LINEN - COTTON OVERSHIRT
Jacket	LINEN - COTTON OVERSHIRT
Jacket	TECHNICAL WINDBREAKER JACKET
Jacket	TECHNICAL WINDBREAKER JACKET
Jacket	TECHNICAL WINDBREAKER JACKET
Jacket	TECHNICAL ANORAK WINDBREAKER
Jacket	TECHNICAL ANORAK WINDBREAKER
Jacket	TECHNICAL ANORAK WINDBREAKER
Jacket	REFLECTIVE EFFECT JACKET
Jacket	REFLECTIVE EFFECT JACKET
Jacket	REFLECTIVE EFFECT JACKET

Rows per page: 20 ▼ 1 – 15 of 15

INDEXING:

Before:

[illegible]

This query has many records since it displays all apparel appropriate for a weather for all users and their subscribed locations. There are thousands of rows in each of these and there is apparel(in the order of 100) info for multiple locations(in the order of 1000) for all users(in the order of 1000s) This is used to show the impact of indices. In our real application we would do it user by user, but it still is a useful scenario to make the best index for our usecase.

Index design 1: apparel(category,product name)

After:

```
create index cat_name_index on apparel (category, product_name);
```

[illegible]

Interpretation: doing an Index on an attribute in the SELECT part of the query doesn't have an impact on the cost of the query. The reason for this is that the select step is done at the end once all the records in the resultset have been fetched. This doesn't decrease the cost of the actual implementation of the query.

Index design 2: apparel(Temp_from):

```
Database Changed
mysql> explain analyze Select a.category, a.product_name
-> From apparel a
-> JOIN( select w.Avg_temp FROM weather w JOIN location l ON w.cityID =l.cityID JOIN usersubscribeslist u ON l.cityID = u.city_ID)
-> AS current_weather
-> Where a.Temp_from <=current_weather.Avg_temp AND a.Temp_to >= current_weather.Avg_temp;
```

```
+-----+
|
```

```
EXPLAIN
```

```
|
```

```
+-----+
|
```

```
-> Nested loop inner join (cost=151650.74 rows=347169) (actual time=0.186..734.045 rows=502004 loops=1)
-> Nested loop inner join (cost=91670.01 rows=100544) (actual time=0.176..353.872 rows=144985 loops=1)
-> Nested loop inner join (cost=1161.85 rows=3027) (actual time=0.097..9.551 rows=3027 loops=1)
-> Covering index scan on l using PRIMARY (cost=102.40 rows=1009) (actual time=0.057..0.660 rows=1009 loops=1)
-> Index lookup on w using cityID (cityID=l.cityID) (cost=0.75 rows=3) (actual time=0.007..0.008 rows=3 loops=1009)
-> Filter: ((a.Temp_from <= w.Avg_temp) and (a.Temp_to >= w.Avg_temp)) (cost=0.00 rows=33) (actual time=0.043..0.110 rows=48 loops=3027)
-> Index range scan on a (re-planned for each iteration) (cost=0.00 rows=299) (actual time=0.011..0.097 rows=85 loops=3027)
-> Covering index lookup on u using city_id (city_id=l.cityID) (cost=0.25 rows=3) (actual time=0.002..0.002 rows=3 loops=144985)
```

Interpretation: Cost decreases massively since the range query used in the Where clause now has an index to efficiently perform the action in. A B-tree hash would allow us to get to the appropriate Temp_from in logn time which is a logarithmic decrease on the time to perform the 'Filter' action as opposed to no index. Then it does a linear search among the leaves that takes as long as the number of leaf nodes matching the condition + 1 (the first one higher than the condition).

Index design 3: apparel(Temp_from,Temp_to):

[illegible]

Interpretation:

Here the values are very similar to the previous one, showing that hashing the Temp_to doesn't really have a massive impact on the increase in performance as having just one of Temp_from and Temp_to in index. A composite index does take up more space and can't be reused as much since a suffix of the index can't be used in other queries to perform them efficiently. So here design 2 seems most efficient and usable.

- b) **print userid , username , gender, cityName of users who subscribe cities that has daily wind speed higher than the avg wind speed of California state in date of 2016-01-03**

```
SELECT u.userid, u.username, u.gender, l.cityName
FROM user u
JOIN usersubscribeslist usl ON u.userid = usl.user_id
JOIN weather w ON usl.city_id = w.cityID AND usl.date = w.date
JOIN location l ON w.cityID = l.cityID
WHERE w.date = '2016-01-03'
AND w.Wind_speed > (
    SELECT AVG(w2.Wind_speed)
    FROM weather w2
    JOIN location l2 ON w2.cityID = l2.cityID
    WHERE l2.State = 'California'
    AND w2.date = '2016-01-03'
)
AND l.State = 'California';
```

userid	username	gender	cityName
965	cmutty	male	El Dorado Hills
611	lupusslugger	female	Davis
981	ThatBlondeWon	female	North Highlands
535	Steve_Mcteta	male	Citrus Heights
765	albertoarchulet	male	Brentwood
83	Beens03	female	Concord
962	TheRealJacoby	male	Ceres
370	Doctor_Math	male	Clovis
78	howugethegirl	female	Bakersfield
850	kolenda23	male	South Whittier
1000	bgdeuce13	male	Rowland Heights
966	GabsterSmitty	female	Cerritos
568	morgane28v	female	Alhambra
865	ShesUntamable	female	Arcadia
953	Haley_1D	female	Azusa

Initial result of default index:

```
| -> Limit: 15 row(s) (cost=169.55 rows=15) (actual time=1.134..1.310 rows=15 loops=1)
| -> Nested loop inner join (cost=169.55 rows=34) (actual time=1.122..1.307 rows=15 loops=1)
| -> Nested loop inner join (cost=157.78 rows=34) (actual time=1.125..1.275 rows=15 loops=1)
| -> Nested loop inner join (cost=137.72 rows=34) (actual time=1.110..1.204 rows=15 loops=1)
| -> Filter: (l.State = 'California') (cost=102.40 rows=101) (actual time=0.292..0.330 rows=33 loops=1)
| -> Table scan on l (cost=102.40 rows=1009) (actual time=0.087..0.277 rows=611 loops=1)
| -> Filter: (w.Wind_speed > (select #2)) (cost=0.25 rows=0.3) (actual time=0.026..0.026 rows=0 loops=33)
| -> Single-row index lookup on w using PRIMARY (date=TIMESTAMP'2016-01-03 00:00:00', cityID=l.cityID) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=33)
| -> Select #2 (subquery in condition; run only once)
| -> Aggregator avg(w2.Wind_speed) (cost=147.81 rows=1) (actual time=0.790..0.790 rows=1 loops=1)
| -> Nested loop inner join (cost=137.72 rows=101) (actual time=0.197..0.769 rows=204 loops=1)
| -> Filter: (l2.State = 'California') (cost=102.40 rows=101) (actual time=0.192..0.376 rows=204 loops=1)
| -> Table scan on l2 (cost=102.40 rows=1009) (actual time=0.023..0.275 rows=1009 loops=1)
| -> Single-row index lookup on w2 using PRIMARY (date=TIMESTAMP'2016-01-03 00:00:00', cityID=l2.cityID) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=204)
| -> Filter: (usl.'date' = TIMESTAMP'2016-01-03 00:00:00') (cost=0.25 rows=1) (actual time=0.004..0.004 rows=1 loops=15)
| -> Covering index lookup on usl using city_id (city_id=l.cityID) (cost=0.25 rows=3) (actual time=0.003..0.004 rows=3 loops=15)
| -> Single-row index lookup on u using PRIMARY (userid=usl.user_id) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=15)
|
```

Index 1 : if create the index on user.username:

```
| -> Limit: 15 row(s) (cost=169.55 rows=15) (actual time=1.129..1.326 rows=15 loops=1)
| -> Nested loop inner join (cost=169.55 rows=34) (actual time=1.127..1.323 rows=15 loops=1)
| -> Nested loop inner join (cost=157.78 rows=34) (actual time=1.118..1.289 rows=15 loops=1)
| -> Nested loop inner join (cost=137.72 rows=34) (actual time=1.100..1.213 rows=15 loops=1)
| -> Filter: (l.State = 'California') (cost=102.40 rows=101) (actual time=0.289..0.341 rows=33 loops=1)
| -> Table scan on l (cost=102.40 rows=1009) (actual time=0.047..0.266 rows=611 loops=1)
| -> Filter: (w.Wind_speed > (select #2)) (cost=0.25 rows=0.3) (actual time=0.026..0.026 rows=0 loops=33)
| -> Single-row index lookup on w using PRIMARY (date=TIMESTAMP'2016-01-03 00:00:00', cityID=l.cityID) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=33)
| -> Select #2 (subquery in condition; run only once)
| -> Aggregator avg(w2.Wind_speed) (cost=147.81 rows=1) (actual time=0.783..0.783 rows=1 loops=1)
| -> Nested loop inner join (cost=137.72 rows=101) (actual time=0.207..0.763 rows=204 loops=1)
| -> Filter: (l2.State = 'California') (cost=102.40 rows=101) (actual time=0.201..0.381 rows=204 loops=1)
| -> Table scan on l2 (cost=102.40 rows=1009) (actual time=0.024..0.279 rows=1009 loops=1)
| -> Single-row index lookup on w2 using PRIMARY (date=TIMESTAMP'2016-01-03 00:00:00', cityID=l2.cityID) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=204)
| -> Filter: (usl.'date' = TIMESTAMP'2016-01-03 00:00:00') (cost=0.25 rows=1) (actual time=0.004..0.005 rows=1 loops=15)
| -> Covering index lookup on usl using city_id (city_id=l.cityID) (cost=0.25 rows=3) (actual time=0.004..0.004 rows=3 loops=15)
| -> Single-row index lookup on u using PRIMARY (userid=usl.user_id) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=15)
|
```

After adding the index, the result remains the same. This is probably caused by our query doesn't use username in order by or where clause and query does not leverage the indexed column in a way that benefits from the index.

Index 2: if create the index on weather.Wind_speed

```
| -> Limit: 15 row(s) (cost=186.68 rows=15) (actual time=0.290..0.492 rows=15 loops=1)
|   -> Nested loop inner join (cost=186.68 rows=52) (actual time=0.290..0.491 rows=15 loops=1)
|     -> Nested loop inner join (cost=160.88 rows=52) (actual time=0.283..0.459 rows=15 loops=1)
|       -> Nested loop inner join (cost=137.72 rows=52) (actual time=0.270..0.391 rows=15 loops=1)
|         -> Filter: (l.State = 'California') (cost=102.40 rows=101) (actual time=0.262..0.300 rows=33 loops=1)
|           -> Table scan on l (cost=102.40 rows=1009) (actual time=0.029..0.242 rows=611 loops=1)
|             -> Filter: (w.Wind_speed > (select #2)) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=0 loops=33)
|               -> Single-row index lookup on w using PRIMARY (date=TIMESTAMP'2016-01-03 00:00:00', cityID=l.cityID) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=33)
|                 -> Select #2 (subquery in condition; run only once)
|                   -> Aggregate: avg(w2.Wind_speed) (cost=147.81 rows=1) (actual time=0.839..0.840 rows=1 loops=1)
|                     -> Nested loop inner join (cost=137.72 rows=101) (actual time=0.282..0.618 rows=204 loops=1)
|                       -> Filter: (l2.State = 'California') (cost=102.40 rows=101) (actual time=0.271..0.449 rows=204 loops=1)
|                         -> Table scan on l2 (cost=102.40 rows=1009) (actual time=0.034..0.330 rows=1009 loops=1)
|                           -> Single-row index lookup on w2 using PRIMARY (date=TIMESTAMP'2016-01-03 00:00:00', cityID=l2.cityID) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=204)
|                             -> Filter: (usl.date = TIMESTAMP'2016-01-03 00:00:00') (cost=0.25 rows=1) (actual time=0.003..0.004 rows=1 loops=15)
|                               -> Covering index lookup on usl using city_id (city_id=l.cityID) (cost=0.25 rows=3) (actual time=0.003..0.004 rows=3 loops=15)
|                                 -> Single-row index lookup on u using PRIMARY (userid=usl.user_id) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=15)
```

After adding an index on wind_speed, the total cost increases. One possible reason for this is that the wind_speed index is not very selective (the threshold (SELECT avg(w2.Wind_speed)) does not filter out too many rows), then scanning the index is probably more costly than doing a full table scan, especially when our result has a lot of rows with wind speeds above the average.

Index design 3: if create the index on user.gender

```
| -> Nested loop inner join (cost=228.45 rows=34) (actual time=1.091..2.030 rows=90 loops=1)
|   -> Nested loop inner join (cost=216.68 rows=34) (actual time=1.084..1.883 rows=90 loops=1)
|     -> Nested loop inner join (cost=188.17 rows=34) (actual time=1.069..1.557 rows=90 loops=1)
|       -> Filter: (l.State = 'California') (cost=102.40 rows=101) (actual time=0.286..0.489 rows=204 loops=1)
|         -> Table scan on l (cost=102.40 rows=1009) (actual time=0.047..0.383 rows=1009 loops=1)
|           -> Filter: (w.Wind_speed > (select #2)) (cost=0.75 rows=0.3) (actual time=0.805..0.805 rows=0 loops=204)
|             -> Single-row index lookup on w using PRIMARY (date=TIMESTAMP'2016-01-03 00:00:00', cityID=l.cityID) (cost=0.75 rows=1) (actual time=0.001..0.001 rows=1 loops=204)
|               -> Select #2 (subquery in condition; run only once)
|                 -> Aggregate: avg(w2.Wind_speed) (cost=198.26 rows=1) (actual time=0.759..0.759 rows=1 loops=1)
|                   -> Nested loop inner join (cost=188.17 rows=101) (actual time=0.257..0.740 rows=204 loops=1)
|                     -> Filter: (l2.State = 'California') (cost=102.40 rows=101) (actual time=0.252..0.450 rows=204 loops=1)
|                       -> Table scan on l2 (cost=102.40 rows=1009) (actual time=0.024..0.341 rows=1009 loops=1)
|                         -> Single-row index lookup on w2 using PRIMARY (date=TIMESTAMP'2016-01-03 00:00:00', cityID=l2.cityID) (cost=0.75 rows=1) (actual time=0.001..0.001 rows=1 loops=204)
|                           -> Filter: (usl.date = TIMESTAMP'2016-01-03 00:00:00') (cost=0.51 rows=1) (actual time=0.002..0.003 rows=1 loops=90)
|                             -> Covering index lookup on usl using city_id (city_id=l.cityID) (cost=0.51 rows=3) (actual time=0.002..0.003 rows=3 loops=90)
|                               -> Single-row index lookup on u using PRIMARY (userid=usl.user_id) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=90)
```

After creating an index on gender, the total cost increases. This is probably because the gender values are not very selective (since many rows have the same values), the database may end up scanning a large portion of the index, which is more expensive than scanning the table directly.

c) Latitude, longitude and cities that have the most extreme weather (highest and lowest temperatures)

```
SELECT l.CityName, l.Latitude, l.Longitude, w.date, w.Max_temp, w.Min_temp
FROM weatherxwardrobe.weather w JOIN weatherxwardrobe.location l ON w.cityID = l.CityID JOIN (
  SELECT cityID, MAX(Max_temp) AS Max_temp, MIN(Min_temp) AS Min_temp
  FROM weatherxwardrobe.weather
  GROUP BY cityID)
extremes ON w.cityID = extremes.cityID AND (w.Max_temp = extremes.Max_temp OR w.Min_temp =
extremes.Min_temp)
ORDER BY l.CityName LIMIT 15
```

CityName	Latitude	Longitude	date	Max_temp	Min_temp
Abilene	32.4543	-99.7384	2016-01-04 00:00:00	55	44
Abilene	32.4543	-99.7384	2016-01-05 00:00:00	46	37
Aguadilla	18.4382	-67.1537	2016-01-05 00:00:00	58	51
Aguadilla	18.4382	-67.1537	2016-01-03 00:00:00	25	12
Akron	41.0798	-81.5219	2016-01-03 00:00:00	71	66
Akron	41.0798	-81.5219	2016-01-04 00:00:00	22	17
Alafaya	28.528	-81.1868	2016-01-04 00:00:00	56	43
Alafaya	28.528	-81.1868	2016-01-03 00:00:00	46	35
Alameda	37.7668	-122.267	2016-01-04 00:00:00	60	50
Alameda	37.7668	-122.267	2016-01-03 00:00:00	41	36
Albany	31.5776	-84.1762	2016-01-05 00:00:00	70	57
Albany	42.6664	-73.7987	2016-01-03 00:00:00	82	76
Albany	42.6664	-73.7987	2016-01-05 00:00:00	23	15
Albany	31.5776	-84.1762	2016-01-04 00:00:00	34	28
Albany	31.5776	-84.1762	2016-01-03 00:00:00	38	28

Rows per page: 20 ▼ 1 – 15 of 15

Before Indexing:

EXPLAIN

```
-> Limit: 16 row(s) (actual time=15.129..15.132 rows=16 loops=1) - ^
> Sort: l.cityName, limit input to 16 row(s) per chunk (actual
time=15.128..15.130 rows=16 loops=1) -> Stream results
(cost=4734.67 rows=1725) (actual time=4.764..14.545 rows=2062
loops=1) -> Nested loop inner join (cost=4734.67 rows=1725)
(actual time=4.751..13.674 rows=2062 loops=1) -> Nested loop
inner join (cost=4130.79 rows=1725) (actual time=4.735..11.818
rows=2062 loops=1) -> Table scan on extremes
(cost=912.11..952.44 rows=3027) (actual time=4.683..4.890
rows=1009 loops=1) -> Materialize (cost=912.10..912.10
rows=3027) (actual time=4.678..4.678 rows=1009 loops=1) ->
Group aggregate: max(weather.Max_temp),
min(weather.Min_temp) (cost=609.40 rows=3027) (actual
time=0.202..4.507 rows=1009 loops=1) -> Index scan on weather
using cityID (cost=306.70 rows=3027) (actual time=0.195..3.918
rows=3027 loops=1) -> Filter: ((w.Max_temp =
extremes.Max_temp) or (w.Min_temp = extremes.Min_temp))
(cost=0.75 rows=1) (actual time=0.006..0.007 rows=2 loops=1009)
-> Index lookup on w using cityID (cityID=extremes.cityID)
(cost=0.75 rows=3) (actual time=0.005..0.006 rows=3 loops=1009)
-> Single-row index lookup on l using PRIMARY
(cityID=extremes.cityID) (cost=0.25 rows=1) (actual
time=0.001..0.001 rows=1 loops=2062)
```

Index 1: CREATE INDEX weather_index ON weather(Max_temp,Min_temp);

```

-> Limit: 16 row(s) (actual time=16.576..16.579 rows=16 loops=1) - ^
> Sort: l.cityName, limit input to 16 row(s) per chunk (actual
time=16.575..16.577 rows=16 loops=1) -> Stream results
(cost=4480.76 rows=1000) (actual time=5.885..16.005 rows=2062
loops=1) -> Nested loop inner join (cost=4480.76 rows=1000)
(actual time=5.876..15.176 rows=2062 loops=1) -> Nested loop
inner join (cost=4130.79 rows=1000) (actual time=5.861..13.201
rows=2062 loops=1) -> Table scan on extremes
(cost=912.11..952.44 rows=3027) (actual time=5.820..6.029
rows=1009 loops=1) -> Materialize (cost=912.10..912.10
rows=3027) (actual time=5.815..5.815 rows=1009 loops=1) ->
Group aggregate: max(weather.Max_temp),
min(weather.Min_temp) (cost=609.40 rows=3027) (actual
time=0.222..5.540 rows=1009 loops=1) -> Index scan on weather
using cityID (cost=306.70 rows=3027) (actual time=0.215..4.803
rows=3027 loops=1) -> Filter: ((w.Max_temp =
extremes.Max_temp) or (w.Min_temp = extremes.Min_temp))
(cost=0.75 rows=0.3) (actual time=0.006..0.007 rows=2
loops=1009) -> Index lookup on w using cityID
(cityID=extremes.cityID) (cost=0.75 rows=3) (actual
time=0.005..0.006 rows=3 loops=1009) -> Single-row index lookup
on l using PRIMARY (cityID=extremes.cityID) (cost=0.25 rows=1)
(actual time=0.001..0.001 rows=1 loops=2062)

```

Indexing the values max_temp and min_temp from the weather table increased cost of query

Index 2:

CREATE INDEX date_index ON weather(date);

```

EXPLAIN
-> Limit: 16 row(s) (actual time=16.315..16.318 rows=16 loops=1) - ^
> Sort: l.cityName, limit input to 16 row(s) per chunk (actual
time=16.314..16.315 rows=16 loops=1) -> Stream results
(cost=4734.67 rows=1725) (actual time=5.492..15.726 rows=2062
loops=1) -> Nested loop inner join (cost=4734.67 rows=1725)
(actual time=5.480..14.882 rows=2062 loops=1) -> Nested loop
inner join (cost=4130.79 rows=1725) (actual time=5.459..12.869
rows=2062 loops=1) -> Table scan on extremes
(cost=912.11..952.44 rows=3027) (actual time=5.395..5.600
rows=1009 loops=1) -> Materialize (cost=912.10..912.10
rows=3027) (actual time=5.390..5.390 rows=1009 loops=1) ->
Group aggregate: max(weather.Max_temp),
min(weather.Min_temp) (cost=609.40 rows=3027) (actual
time=0.218..5.213 rows=1009 loops=1) -> Index scan on weather
using cityID (cost=306.70 rows=3027) (actual time=0.211..4.583
rows=3027 loops=1) -> Filter: ((w.Max_temp =
extremes.Max_temp) or (w.Min_temp = extremes.Min_temp))
(cost=0.75 rows=1) (actual time=0.006..0.007 rows=2 loops=1009)
-> Index lookup on w using cityID (cityID=extremes.cityID)
(cost=0.75 rows=3) (actual time=0.006..0.006 rows=3 loops=1009)
-> Single-row index lookup on l using PRIMARY
(cityID=extremes.cityID) (cost=0.25 rows=1) (actual
time=0.001..0.001 rows=1 loops=2062)

```

After indexing the date values from the weather table, the cost remains the same. This tells us that it takes the same amount of effort to perform the query.

Index 3:

CREATE INDEX latlong_index ON location(latitude, longitude);

EXPLAIN

```
-> Limit: 16 row(s) (actual time=15.506..15.508 rows=16 loops=1) - ^
> Sort: l.cityName, limit input to 16 row(s) per chunk (actual
time=15.505..15.507 rows=16 loops=1) -> Stream results
(cost=4734.67 rows=1725) (actual time=5.063..14.945 rows=2062
loops=1) -> Nested loop inner join (cost=4734.67 rows=1725)
(actual time=5.053..14.108 rows=2062 loops=1) -> Nested loop
inner join (cost=4130.79 rows=1725) (actual time=5.040..12.053
rows=2062 loops=1) -> Table scan on extremes
(cost=912.11..952.44 rows=3027) (actual time=5.011..5.219
rows=1009 loops=1) -> Materialize (cost=912.10..912.10
rows=3027) (actual time=5.007..5.007 rows=1009 loops=1) ->
Group aggregate: max(weather.Max_temp),
min(weather.Min_temp) (cost=609.40 rows=3027) (actual
time=0.192..4.838 rows=1009 loops=1) -> Index scan on weather
using cityID (cost=306.70 rows=3027) (actual time=0.186..4.210
rows=3027 loops=1) -> Filter: ((w.Max_temp =
extremes.Max_temp) or (w.Min_temp = extremes.Min_temp))
(cost=0.75 rows=1) (actual time=0.006..0.007 rows=2 loops=1009)
-> Index lookup on w using cityID (cityID=extremes.cityID)
(cost=0.75 rows=3) (actual time=0.005..0.006 rows=3 loops=1009)
-> Single-row index lookup on l using PRIMARY
(cityID=extremes.cityID) (cost=0.25 rows=1) (actual
time=0.001..0.001 rows=1 loops=2062)
```

After indexing the latitude and longitude values from the location table, the cost remains the same. This tells us that it takes the same amount of effort to perform the query.

Index 4:

CREATE INDEX city_index ON location(cityName);

EXPLAIN

```
-> Limit: 16 row(s) (actual time=15.607..15.609 rows=16 loops=1) - ^
> Sort: l.cityName, limit input to 16 row(s) per chunk (actual
time=15.606..15.607 rows=16 loops=1) -> Stream results
(cost=4734.67 rows=1725) (actual time=5.023..15.012 rows=2062
loops=1) -> Nested loop inner join (cost=4734.67 rows=1725)
(actual time=5.013..14.138 rows=2062 loops=1) -> Nested loop
inner join (cost=4130.79 rows=1725) (actual time=4.999..12.136
rows=2062 loops=1) -> Table scan on extremes
(cost=912.11..952.44 rows=3027) (actual time=4.969..5.182
rows=1009 loops=1) -> Materialize (cost=912.10..912.10
rows=3027) (actual time=4.965..4.965 rows=1009 loops=1) ->
Group aggregate: max(weather.Max_temp),
min(weather.Min_temp) (cost=609.40 rows=3027) (actual
time=0.193..4.788 rows=1009 loops=1) -> Index scan on weather
using cityID (cost=306.70 rows=3027) (actual time=0.187..4.212
rows=3027 loops=1) -> Filter: ((w.Max_temp =
extremes.Max_temp) or (w.Min_temp = extremes.Min_temp))
(cost=0.75 rows=1) (actual time=0.006..0.007 rows=2 loops=1009)
-> Index lookup on w using cityID (cityID=extremes.cityID)
(cost=0.75 rows=3) (actual time=0.005..0.006 rows=3 loops=1009)
-> Single-row index lookup on l using PRIMARY
(cityID=extremes.cityID) (cost=0.25 rows=1) (actual
time=0.001..0.001 rows=1 loops=2062)
```

After indexing the cityName values from location, the cost to perform the query remains the same.

d) Cities with temperature higher than the Average temperatures of all over a date range.

EXPLAIN ANALYZE QUERY FOR INDEXING:

```
EXPLAIN ANALYZE SELECT l.CityName, w.date, w.Max_temp, w.Wind_speed, w.Category
FROM weatherxwardrobe.weather w
JOIN weatherxwardrobe.location l ON w.cityID = l.CityID
WHERE w.Max_temp > (
    SELECT AVG(Max_temp)
    FROM weatherxwardrobe.weather
    WHERE date BETWEEN '2016-01-03 00:00:00' AND '2016-01-05 00:00:00'
)
AND w.date IN (
    SELECT DISTINCT date
    FROM weatherxwardrobe.weather
    WHERE Max_temp > 30
    GROUP BY date
    HAVING COUNT(*) > 2
)
ORDER BY w.date, l.CityName
LIMIT 15;
```

DEFAULT INDEXED RESULT:

-> Limit: 15 row(s) (actual time=13.062..13.065 rows=15 loops=1) -  Toggle f

> Sort: w.`date`, l.cityName, limit input to 15 row(s) per chunk
 (actual time=13.061..13.063 rows=15 loops=1) -> Stream results
 (cost=458.00 rows=1009) (actual time=7.494..12.553 rows=1493 loops=1) -> Nested loop inner join (cost=458.00 rows=1009)
 (actual time=7.487..11.971 rows=1493 loops=1) -> Filter:
 ((w.Max_temp > (select #2)) and <in_optimizer>(w.`date`,w.`date` in
 (select #3))) (cost=104.89 rows=1009) (actual time=7.459..9.713
 rows=1493 loops=1) -> Table scan on w (cost=104.89 rows=3027)
 (actual time=0.071..1.476 rows=3027 loops=1) -> Select #2
 (subquery in condition; run only once) -> Aggregate:
 avg(weather.Max_temp) (cost=454.76 rows=1) (actual
 time=5.887..5.888 rows=1 loops=1) -> Filter: (weather.`date`
 between '2016-01-03 00:00:00' and '2016-01-05 00:00:00')
 (cost=303.46 rows=1513) (actual time=0.049..5.288 rows=3027
 loops=1) -> Index range scan on weather using PRIMARY over
 ('2016-01-03 00:00:00' <= date <= '2016-01-05 00:00:00')
 (cost=303.46 rows=1513) (actual time=0.038..2.114 rows=3027
 loops=1) -> Select #3 (subquery in condition; run only once) ->
 Filter: ((w.`date` = `<materialized_subquery>`.`date`))
 (cost=508.58..508.58 rows=1) (actual time=0.362..0.362 rows=1
 loops=4) -> Limit: 1 row(s) (cost=508.48..508.48 rows=1) (actual
 time=0.357..0.357 rows=1 loops=4) -> Index lookup on
 <materialized_subquery> using <auto_distinct_key> (date=w.`date`)
 (actual time=0.356..0.356 rows=1 loops=4) -> Materialize with
 deduplication (cost=508.48..508.48 rows=1009) (actual
 time=1.405..1.405 rows=3 loops=1) -> Filter: (count(0) > 2)
 (cost=407.59 rows=1009) (actual time=0.485..1.384 rows=3
 loops=1) -> Group aggregate: count(0) (cost=407.59 rows=1009)
 (actual time=0.484..1.381 rows=3 loops=1) -> Filter:
 (weather.Max_temp > 30) (cost=306.70 rows=1009) (actual
 time=0.043..1.210 rows=2564 loops=1) -> Index scan on weather
 using PRIMARY (cost=306.70 rows=3027) (actual
 time=0.040..0.938 rows=3027 loops=1) -> Single-row index lookup
 on l using PRIMARY (cityID=w.cityID) (cost=0.25 rows=1) (actual
 time=0.001..0.001 rows=1 loops=1493)

FIRST INDEXING METHOD:

CREATE INDEX weather_index ON weather(Max_temp,Min_temp);

RESULT:

-> Limit: 15 row(s) (actual time=6.611..6.613 rows=15 loops=1) ->  Toggle
 Sort: w.`date`, l.cityName, limit input to 15 row(s) per chunk (actual time=6.610..6.611 rows=15 loops=1) -> Stream results
 (cost=824.95 rows=1493) (actual time=2.064..6.166 rows=1493 loops=1) -> Nested loop inner join (cost=824.95 rows=1493) (actual time=2.059..5.647 rows=1493 loops=1) -> Filter: ((w.Max_temp > (select #2)) and <in_optimizer>(w.`date`,w.`date` in (select #3))) (cost=302.40 rows=1493) (actual time=2.044..3.612 rows=1493 loops=1) -> Covering index range scan on w using weather_index over (47 < Max_temp) (cost=302.40 rows=1493) (actual time=0.022..0.521 rows=1493 loops=1) -> Select #2 (subquery in condition; run only once) -> Aggregate: avg(weather.Max_temp) (cost=454.76 rows=1) (actual time=3.358..3.359 rows=1 loops=1) -> Filter: (weather.`date` between '2016-01-03 00:00:00' and '2016-01-05 00:00:00') (cost=303.46 rows=1513) (actual time=0.033..2.980 rows=3027 loops=1) -> Index range scan on weather using PRIMARY over ('2016-01-03 00:00:00' <= date <= '2016-01-05 00:00:00') (cost=303.46 rows=1513) (actual time=0.028..1.344 rows=3027 loops=1) -> Select #3 (subquery in condition; run only once) -> Filter: ((w.`date` = <materialized_subquery>.`date`)) (cost=0.00..0.00 rows=0) (actual time=0.004..0.004 rows=1 loops=652) -> Limit: 1 row(s) (cost=0.00..0.00 rows=0) (actual time=0.004..0.004 rows=1 loops=652) -> Index lookup on <materialized_subquery> using <auto_distinct_key> (date=w.`date`) (actual time=0.003..0.003 rows=1 loops=652) -> Materialize with deduplication (cost=0.00..0.00 rows=0) (actual time=2.009..2.009 rows=3 loops=1) -> Filter: (count(0) > 2) (actual time=2.002..2.003 rows=3 loops=1) -> Table scan on <temporary> (actual time=2.000..2.001 rows=3 loops=1) -> Aggregate using temporary table (actual time=1.998..1.998 rows=3 loops=1) -> Filter: (weather.Max_temp > 30) (cost=518.70 rows=2564) (actual time=0.010..1.018 rows=2564 loops=1) -> Covering index range scan on weather using weather_index over (30 < Max_temp) (cost=518.70 rows=2564) (actual time=0.009..0.810 rows=2564 loops=1) -> Single-row index lookup on l using PRIMARY (cityID=w.cityID) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=1493)

Increases the cost for the query. Meaning, by indexing the Max_temp and Min_temp values from the weather table, it takes more effort to perform the query.

DELETED INDEX

ALTER TABLE weather DROP INDEX weather_index;

Command used to delete the indexing method.

SECOND INDEXING METHOD:

CREATE INDEX location_index ON location(cityID);

Result:

```

-> Limit: 15 row(s) (actual time=8.812..8.815 rows=15 loops=1) -> ^
Sort: w.`date`, l.cityName, limit input to 15 row(s) per chunk (actual
time=8.811..8.813 rows=15 loops=1) -> Stream results
(aggers (cost=458.00 rows=1009) (actual time=4.072..8.390 rows=1493
loops=1) -> Nested loop inner join (cost=458.00 rows=1009)
(actual time=4.067..7.865 rows=1493 loops=1) -> Filter:
((w.Max_temp > (select #2)) and <in_optimizer>(w.`date`,w.`date` in
(select #3))) (cost=104.89 rows=1009) (actual time=4.051..5.822
rows=1493 loops=1) -> Table scan on w (cost=104.89 rows=3027)
(actual time=0.032..0.973 rows=3027 loops=1) -> Select #2
(subquery in condition; run only once) -> Aggregate:
avg(weather.Max_temp) (cost=454.76 rows=1) (actual
time=2.756..2.756 rows=1 loops=1) -> Filter: (weather.`date`
between '2016-01-03 00:00:00' and '2016-01-05 00:00:00')
(cost=303.46 rows=1513) (actual time=0.009..2.470 rows=3027
loops=1) -> Index range scan on weather using PRIMARY over
('2016-01-03 00:00:00' <= date <= '2016-01-05 00:00:00')
(cost=303.46 rows=1513) (actual time=0.007..1.085 rows=3027
loops=1) -> Select #3 (subquery in condition; run only once) ->
Filter: ((w.`date` = `<materialized_subquery>`.`date`))
(cost=508.58..508.58 rows=1) (actual time=0.312..0.312 rows=1
loops=4) -> Limit: 1 row(s) (cost=508.48..508.48 rows=1) (actual
time=0.311..0.311 rows=1 loops=4) -> Index lookup on
<materialized_subquery> using <auto_distinct_key> (date=w.`date`)
(actual time=0.311..0.311 rows=1 loops=4) -> Materialize with
deduplication (cost=508.48..508.48 rows=1009) (actual
time=1.235..1.235 rows=3 loops=1) -> Filter: (count(0) > 2)
(cost=407.59 rows=1009) (actual time=0.419..1.226 rows=3
loops=1) -> Group aggregate: count(0) (cost=407.59 rows=1009)
(actual time=0.417..1.224 rows=3 loops=1) -> Filter:
(weather.Max_temp > 30) (cost=306.70 rows=1009) (actual
time=0.028..1.072 rows=2564 loops=1) -> Index scan on weather
using PRIMARY (cost=306.70 rows=3027) (actual
time=0.026..0.833 rows=3027 loops=1) -> Single-row index lookup
on l using PRIMARY (cityID=w.cityID) (cost=0.25 rows=1) (actual
time=0.001..0.001 rows=1 loops=1493)

```

The cost is not changed by indexing the cityID value from the location table. This could mean that it does not change the performance to index the values from the location table for this query.

DELETED INDEX:

ALTER TABLE location DROP INDEX location_index;

Deleted second indexing method.

THIRD INDEXING METHOD:

CREATE INDEX final_index ON weather(Category,Wind_speed);

Result:

-> Limit: 15 row(s) (actual time=8.812..8.815 rows=15 loops=1) -> ^
 Sort: w.`date`, l.cityName, limit input to 15 row(s) per chunk (actual
 time=8.811..8.813 rows=15 loops=1) -> Stream results
 (cost=458.00 rows=1009) (actual time=4.072..8.390 rows=1493
 loops=1) -> Nested loop inner join (cost=458.00 rows=1009)
 (actual time=4.067..7.865 rows=1493 loops=1) -> Filter:
 ((w.Max_temp > (select #2)) and <in_optimizer>(w.`date`,w.`date` in
 (select #3))) (cost=104.89 rows=1009) (actual time=4.051..5.822
 rows=1493 loops=1) -> Table scan on w (cost=104.89 rows=3027)
 (actual time=0.032..0.973 rows=3027 loops=1) -> Select #2
 (subquery in condition; run only once) -> Aggregate:
 avg(weather.Max_temp) (cost=454.76 rows=1) (actual
 time=2.756..2.756 rows=1 loops=1) -> Filter: (weather.`date`
 between '2016-01-03 00:00:00' and '2016-01-05 00:00:00')
 (cost=303.46 rows=1513) (actual time=0.009..2.470 rows=3027
 loops=1) -> Index range scan on weather using PRIMARY over
 ('2016-01-03 00:00:00' <= date <= '2016-01-05 00:00:00')
 (cost=303.46 rows=1513) (actual time=0.007..1.085 rows=3027
 loops=1) -> Select #3 (subquery in condition; run only once) ->
 Filter: ((w.`date` = `<materialized_subquery>`.`date`))
 (cost=508.58..508.58 rows=1) (actual time=0.312..0.312 rows=1
 loops=4) -> Limit: 1 row(s) (cost=508.48..508.48 rows=1) (actual
 time=0.311..0.311 rows=1 loops=4) -> Index lookup on
 <materialized_subquery> using <auto_distinct_key> (date=w.`date`)
 (actual time=0.311..0.311 rows=1 loops=4) -> Materialize with
 deduplication (cost=508.48..508.48 rows=1009) (actual
 time=1.235..1.235 rows=3 loops=1) -> Filter: (count(0) > 2)
 (cost=407.59 rows=1009) (actual time=0.419..1.226 rows=3
 loops=1) -> Group aggregate: count(0) (cost=407.59 rows=1009)
 (actual time=0.417..1.224 rows=3 loops=1) -> Filter:
 (weather.Max_temp > 30) (cost=306.70 rows=1009) (actual
 time=0.028..1.072 rows=2564 loops=1) -> Index scan on weather
 using PRIMARY (cost=306.70 rows=3027) (actual
 time=0.026..0.833 rows=3027 loops=1) -> Single-row index lookup
 on l using PRIMARY (cityID=w.cityID) (cost=0.25 rows=1) (actual
 time=0.001..0.001 rows=1 loops=1493)

Didn't change anything. This means that indexing the category and wind_speed data values from the weather table did not increase or decrease the efficiency of the query.

DELETED INDEXING:

ALTER TABLE weather DROP INDEX final_index;

Deleted the third indexing method.

D QUERY:

```
SELECT l.CityName, w.date, w.Max_temp
FROM weatherxwardrobe.weather w
JOIN weatherxwardrobe.location l ON w.cityID = l.CityID
WHERE w.Max_temp > (
  SELECT AVG(Max_temp)
  FROM weatherxwardrobe.weather
  WHERE date BETWEEN '2016-01-03 00:00:00' AND '2016-01-05 00:00:00'
```

CityName	date	Max_temp
Akron	2016-01-03 00:00:00	71
Albany	2016-01-03 00:00:00	82
Albuquerque	2016-01-03 00:00:00	51
Alexandria	2016-01-03 00:00:00	52
Allen	2016-01-03 00:00:00	48
Aloha	2016-01-03 00:00:00	54
Alton	2016-01-03 00:00:00	63
Altoona	2016-01-03 00:00:00	55
Ankeny	2016-01-03 00:00:00	52
Anniston	2016-01-03 00:00:00	51
Antelope	2016-01-03 00:00:00	61
Apex	2016-01-03 00:00:00	62
Apopka	2016-01-03 00:00:00	52
Arvada	2016-01-03 00:00:00	54
Auburn	2016-01-03 00:00:00	49
Austin	2016-01-03 00:00:00	60

D QUERY RESULTS (TOP 15)

Changes for stage 2:

1.ER diagram remade completely:

Weather is a weak entity with a supporting relationship with Location since each weather is uniquely identified by the location of the weather. For eg. "Snowy" isn't unique since multiple locations can have the same snowy weather. So "Snowy","Urbana" is a unique combination.

Made user-list into a many to many relationship between user and weather, thus the schema remains the same but the relationship is more clearly defined.

Added more information to Location entity set.Now State is uniquely defined by City_Id and weather is uniquely defined by City_id and date.

2. Fixed inconsistencies between schema and ER diagram.